

# Benchmark Report — California 2021 StationMax Uniform Kernel

**Date:** 2026-03-13

**Dataset:** California 2021

**Risk map:** Pyrologix ignition probability (static, trained on 2006–2020, no data leakage for 2021)

---

## Overview

This benchmark tracks the new **StationMax uniform-kernel** placement strategy. It keeps the same high-level formulation as the greedy-kernel StationMax model:

- coverage from **different charging stations does not add** at a cell
- each cell only receives credit from its **best selected station**
- candidate station-cell pairs are restricted to the masked **one-way reachable** zone with  $\text{floor}(\text{max\_battery\_time} / 2) = 3$  operational cells

The difference is the **same-station multi-drone kernel**:

- for a station with  $k$  drones and  $R$  feasible masked reachable cells
- each reachable cell gets coverage  $\min(1, k * B / R)$
- where  $B = \text{max\_battery\_time} = 7$  operational substeps

So when all  $7 * 7 = 49$  cells are feasible, this reduces to  $\min(1, k / 7)$ .

For budgets above 300M, the objective also includes the small regularization term  $-0.1 * \text{budget\_used}$ , which prefers cheaper solutions once coverage is already saturated.

---

## Common Setup

| Parameter                     | Value                                          |
|-------------------------------|------------------------------------------------|
| Data grid                     | 1309 × 805 (~1 km/cell)                        |
| Operational grid              | 261 × 161 (5 km/cell)                          |
| Drone speed                   | 600 m/min                                      |
| Coverage radius               | 2900 m                                         |
| Battery                       | 1 h = 7 operational substeps                   |
| One-way reach used in pruning | 3 operational cells                            |
| Candidate pre-filtering       | top 20% by risk / coverage potential           |
| Solver                        | Gurobi                                         |
| Time limit per run            | 600 s default; 1800 s for the final 100M rerun |

---

## Fire Locations

---

## 500M Run

The 500M case was executed after introducing the new StationMax uniform strategy.

Observed result:

| Metric            | Value          |
|-------------------|----------------|
| Termination       | TIME_LIMIT     |
| Preprocessing     | 2.37 s         |
| Model creation    | 6.30 s         |
| Solve time        | 604.79 s       |
| Ground sensors    | 0              |
| Charging stations | 0              |
| Drones            | 0              |
| Budget used       | 0.00M / 500.0M |

The solver returned objective -0.0 with a positive bound, so this should be interpreted as **failure to find a meaningful feasible incumbent within the 10-minute cap**, not as a valid empty optimum.

## Plots

---

## Reproduction

To rerun the full benchmark set:

`bash report/benchmark_2021_stationmax_uniform_kernel/reproduce_benchmark_2021_stationmax_un`




Figure 1: California 2021 fire locations

**Pyrologix + cluster zones + fires [GaussianBudget500M\_StationMaxUniform\_261x161\_mean]**  
**0 clusters · 931 fires (0 discoverable, 931 non-disc)**




Figure 2: 500M StationMax uniform kernel — data scale

Pyrologix (opt scale) + cluster zones + fires [GaussianBudget500M\_StationMaxUniform\_261x161\_mean]  
 0 clusters · 931 fires (grid 261×161, 5 km / cell)




Figure 3: 500M StationMax uniform kernel — operational scale